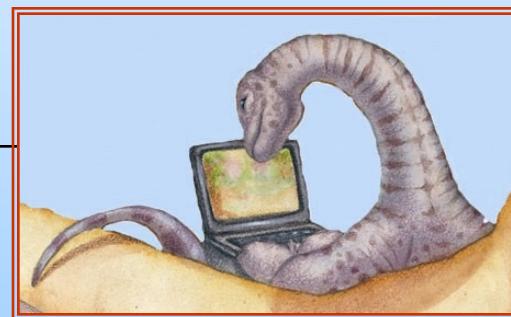




The picture can't be displayed.

# Chapter 2: Operating-System Structures





# Chapter 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- System Calls
- Types of System Calls
- System Programs
- Operating System Design and Implementation
- Operating System Structure
- Virtual Machines
- Operating System Generation
- System Boot





# Objectives

- To describe the **services** an operating system provides to users, processes, and other systems
- To discuss the various ways of **structuring** an operating system
- To explain how operating systems are installed and customized and how they boot





# Chapter 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- System Calls
- Types of System Calls
- System Programs
- Operating System Design and Implementation
- Operating System Structure
- Virtual Machines
- Operating System Generation
- System Boot





# Operating System Services

These are from  
the user's  
perspective

- One set of operating-system services provides functions that are helpful to the user:

- **User interface** - Almost all operating systems have a user interface (UI)
  - ▶ Varies between Command-Line (CLI), Graphics User Interface (GUI), Batch
- **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)
- **I/O operations** - A running program may require I/O, which may involve a file or an I/O device.
- **File-system manipulation** - The file system is of particular interest. Obviously, programs need to read and write files and directories, create and delete them, search them, list file information, permission management.





# Operating System Services (Cont.)

These are from  
the user's  
perspective

- One set of operating-system services provides functions that are helpful to the user (Cont):

- **Communications** – Processes may exchange information, on the same computer or between computers over a network
  - ▶ Communications may be via shared memory or through message passing (packets moved by the OS)
- **Error detection** – OS needs to be constantly aware of possible errors
  - ▶ May occur in the CPU and memory hardware, in I/O devices, in user program
  - ▶ For each type of error, OS should take the appropriate action to ensure correct and consistent computing
  - ▶ Debugging facilities can greatly enhance the user's and programmer's abilities to efficiently use the system





# Chapter 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- System Calls
- Types of System Calls
- System Programs
- Operating System Design and Implementation
- Operating System Structure
- Virtual Machines
- Operating System Generation
- System Boot





# User Operating System Interface - CLI

CLI (Command Line Interface) allows direct command entry

- ▶ Sometimes implemented in **kernel**, sometimes by **systems program**
- ▶ Sometimes multiple flavors implemented – **shells**
- ▶ Primarily fetches a command from user and executes it
  - Sometimes commands built-in (**DOS**), sometimes just names of programs (**Unix**)
    - » If the latter, adding new features doesn't require shell modification





# User Operating System Interface - GUI

- User-friendly **desktop** metaphor interface
  - Usually mouse, keyboard, and monitor
  - **Icons** represent files, programs, actions, etc
  - Various mouse buttons over objects in the interface cause various actions (provide information, options, execute function, open directory (known as a **folder**)
  - Invented at Xerox PARC
- Many systems now include both CLI and GUI interfaces
  - Microsoft Windows is GUI with CLI “command” shell
  - Apple Mac OS X as “Aqua” GUI interface with UNIX kernel underneath and shells available
  - Solaris is CLI with optional GUI interfaces (Java Desktop, KDE)





# Chapter 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- **System Calls**
- Types of System Calls
- System Programs
- Operating System Design and Implementation
- Operating System Structure
- Virtual Machines
- Operating System Generation
- System Boot





# System Calls

- Programming interface to the services provided by the OS
- Typically written in a high-level language (C or C++)
- Mostly accessed by programs via a high-level **Application Program Interface (API)** rather than direct system call use
- Three most common APIs are Win32 API for Windows, POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X), and Java API for the Java virtual machine (JVM)
- Why use APIs rather than system calls?  
*We'll see.*





# System Calls – How do they look like?

- They look like following in **header files**:

```
#include <sys/socket.h>

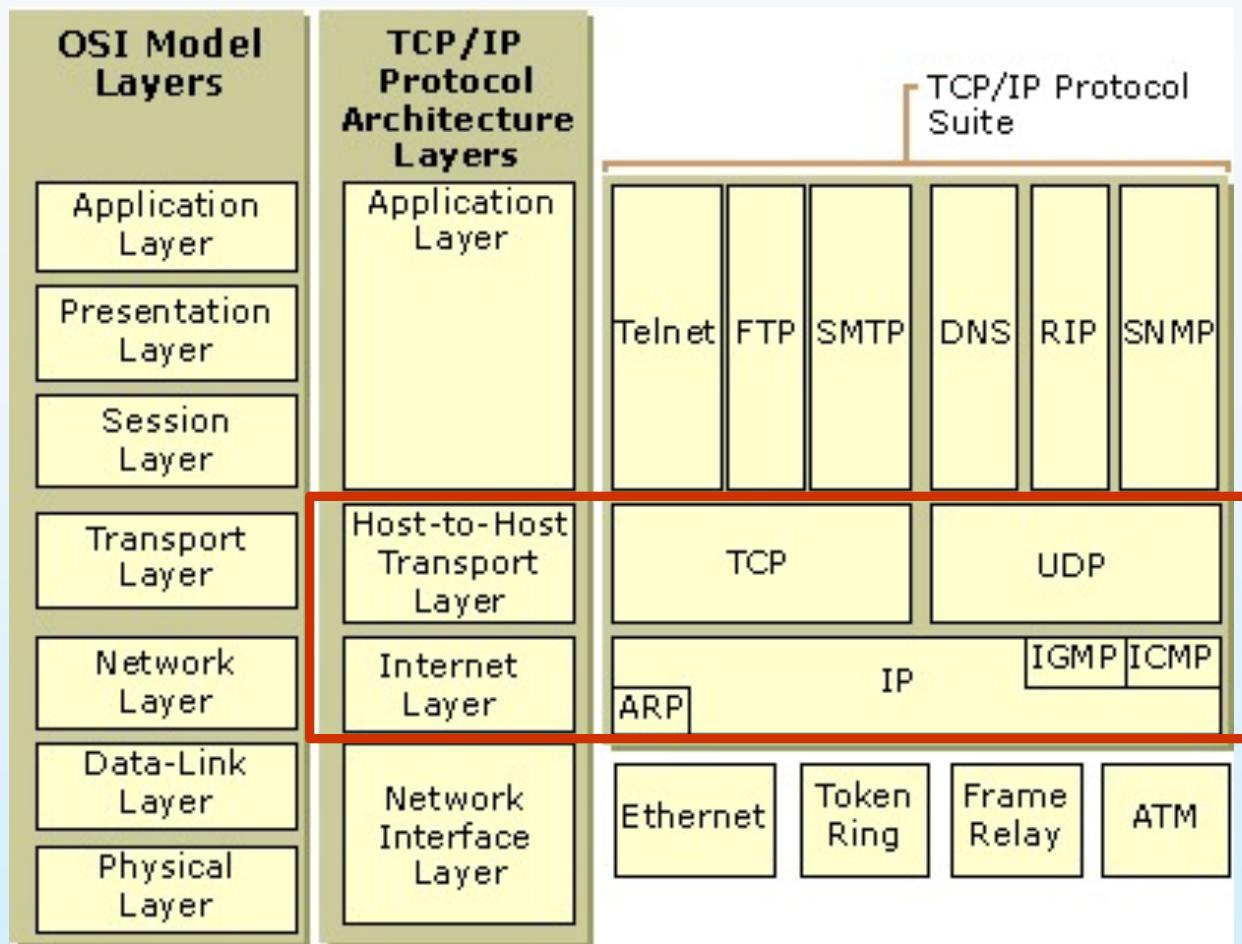
int socket(int socket_family, int socket_type, int protocol);
int bind(int sockfd, const struct sockaddr *addr, socklen_t
addrlen);
int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
int connect(int sockfd, const struct sockaddr *addr,
socklen_t addrlen);
```





# System Calls –How do they look like?

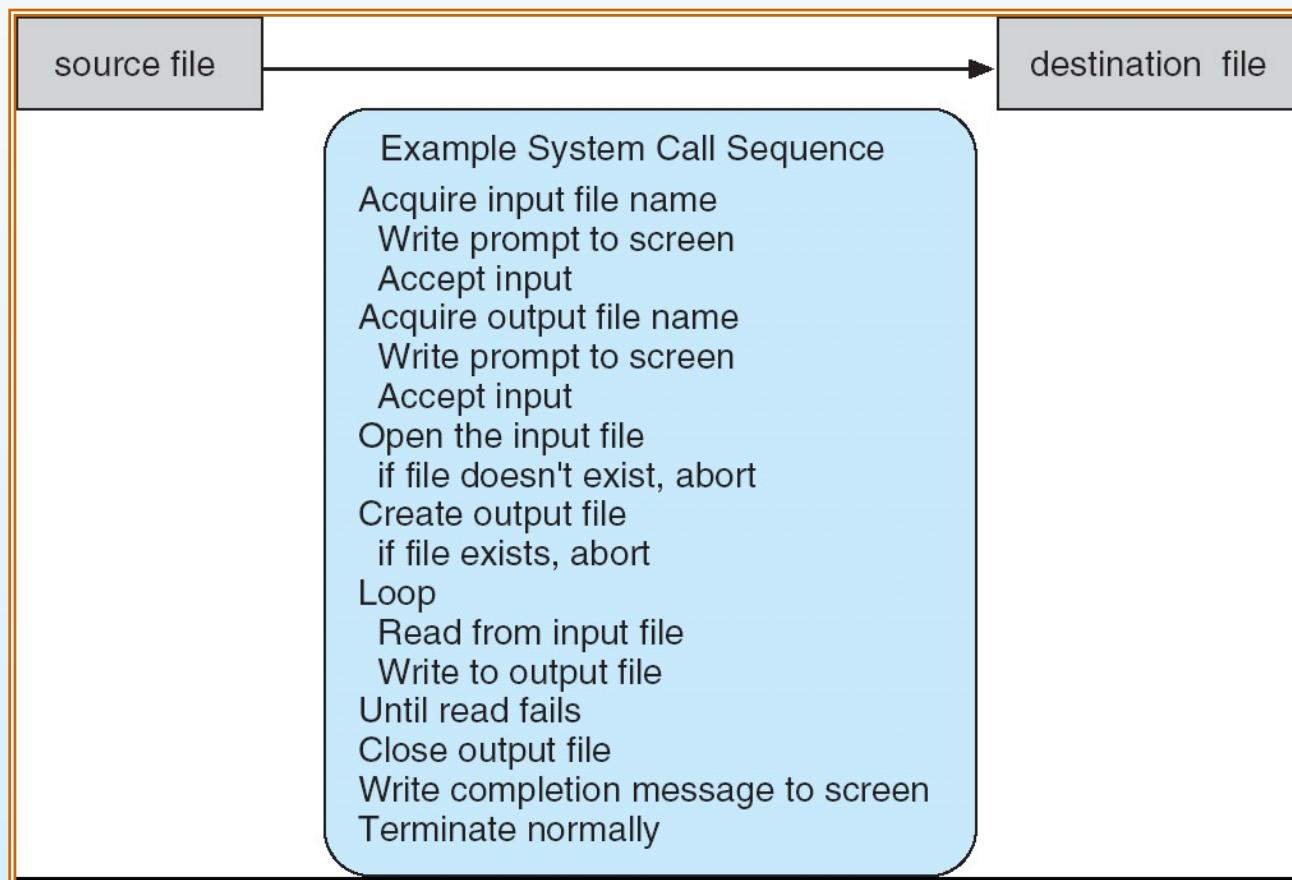
They look like following in **layered model**.





# Example of System Calls

- System call sequence to copy the contents of one file to another file





# Example of Standard API

- Consider the ReadFile() function in the Win32 API—a function for reading from a file

```
return value
      ↓
BOOL ReadFile c (HANDLE file,
                  LPVOID buffer,
                  DWORD bytes To Read,
                  LPDWORD bytes Read,
                  LPOVERLAPPED ovl);
                  ↑
function name
```

A description of the parameters passed to ReadFile()

The diagram shows the signature of the ReadFile() function. The return value is a BOOL type. The function name is ReadFile. It takes five parameters: file (HANDLE), buffer (LPVOID), bytesToRead (DWORD), bytesRead (LPDWORD), and ovl (LPOVERLAPPED). The parameters bytesToRead, bytesRead, and ovl are grouped together with a bracket on the right.

- A description of the parameters passed to ReadFile()
  - HANDLE file—the file to be read
  - LPVOID buffer—a buffer where the data will be read into and written from
  - DWORD bytesToRead—the number of bytes to be read into the buffer
  - LPDWORD bytesRead—the number of bytes read during the last read
  - LPOVERLAPPED ovl—indicates if overlapped I/O is being used





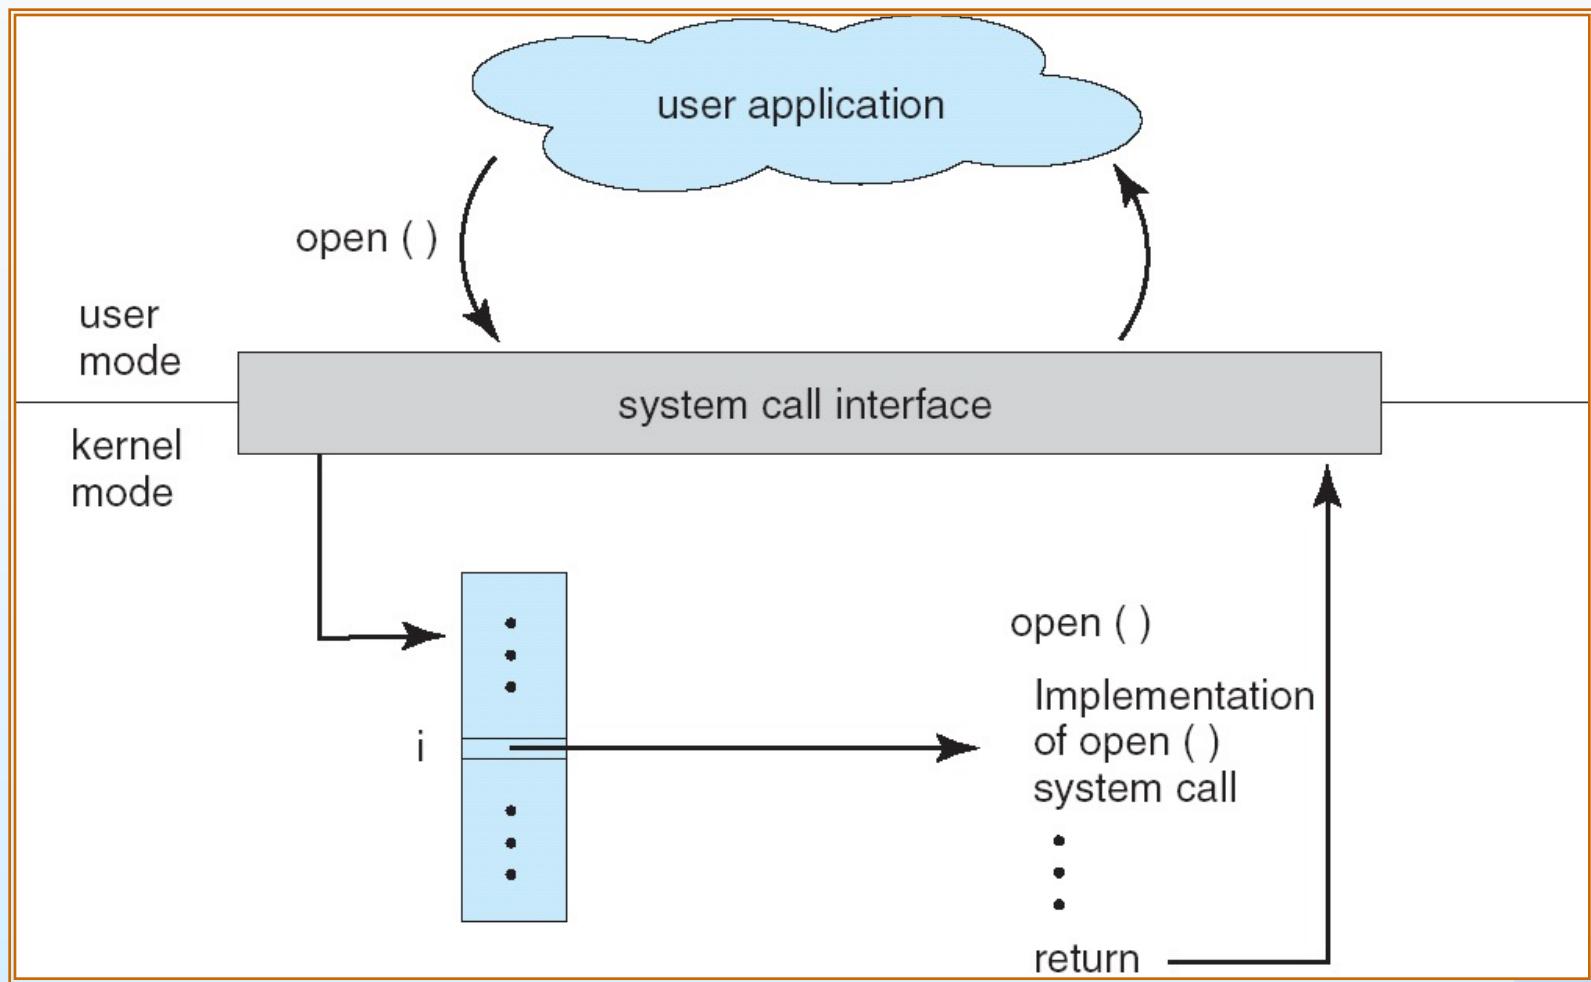
# System Call Implementation

- Typically, a **number** associated with each system call
  - System-call interface maintains a **table** indexed according to these numbers
- The system call interface invokes intended system call in OS kernel and returns status of the system call and any return values
- The caller needs to know nothing about how the system call is implemented
  - Just needs to obey API and understand what OS will do as a result call
  - Most **details** of OS interface hidden from programmer by API
    - ▶ Managed by run-time support library (set of functions built into libraries included with compiler)





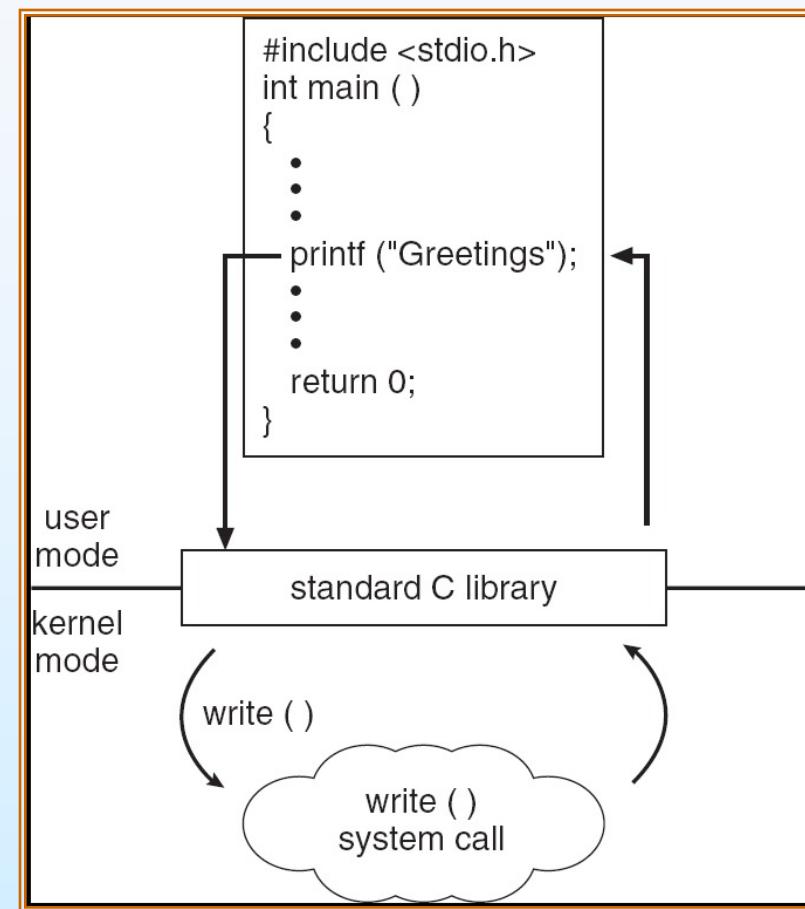
# API – System Call – OS Relationship





# Standard C Library Example

- C program invoking printf() library call, which calls write() system call





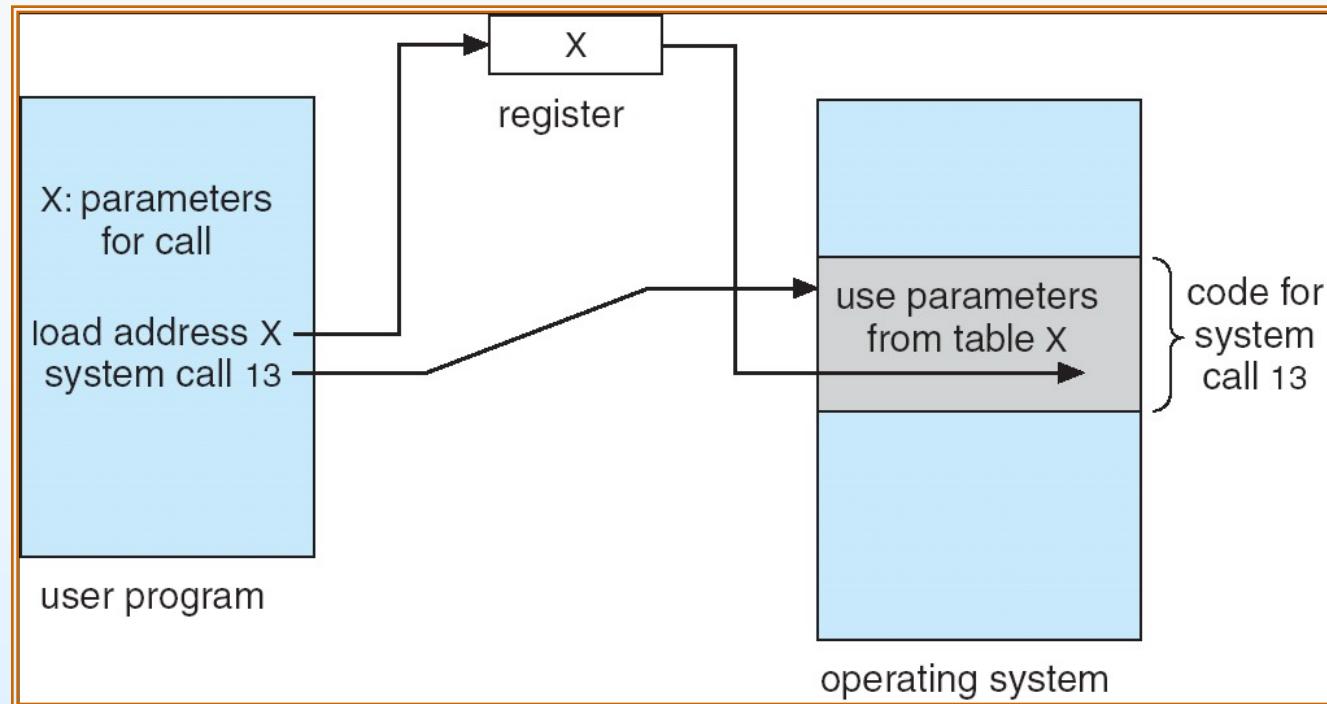
# System Call Parameter Passing

- Often, more information is required than simply identity of desired system call
  - Exact type and amount of information vary according to OS and call
- Three general methods used to pass parameters to the OS
  - Simplest: pass the parameters in *registers*
    - ▶ In some cases, may be more parameters than registers
  - Parameters stored in a *block*, or table, in memory, and address of block passed as a parameter in a register
    - ▶ This approach taken by Linux and Solaris
  - Parameters placed, or *pushed*, onto the *stack* by the program and *popped* off the stack by the operating system
  - Block and stack methods do not limit the number or length of parameters being passed





# Parameter Passing via Table





# Chapter 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- System Calls
- **Types of System Calls**
- System Programs
- Operating System Design and Implementation
- Operating System Structure
- Virtual Machines
- Operating System Generation
- System Boot





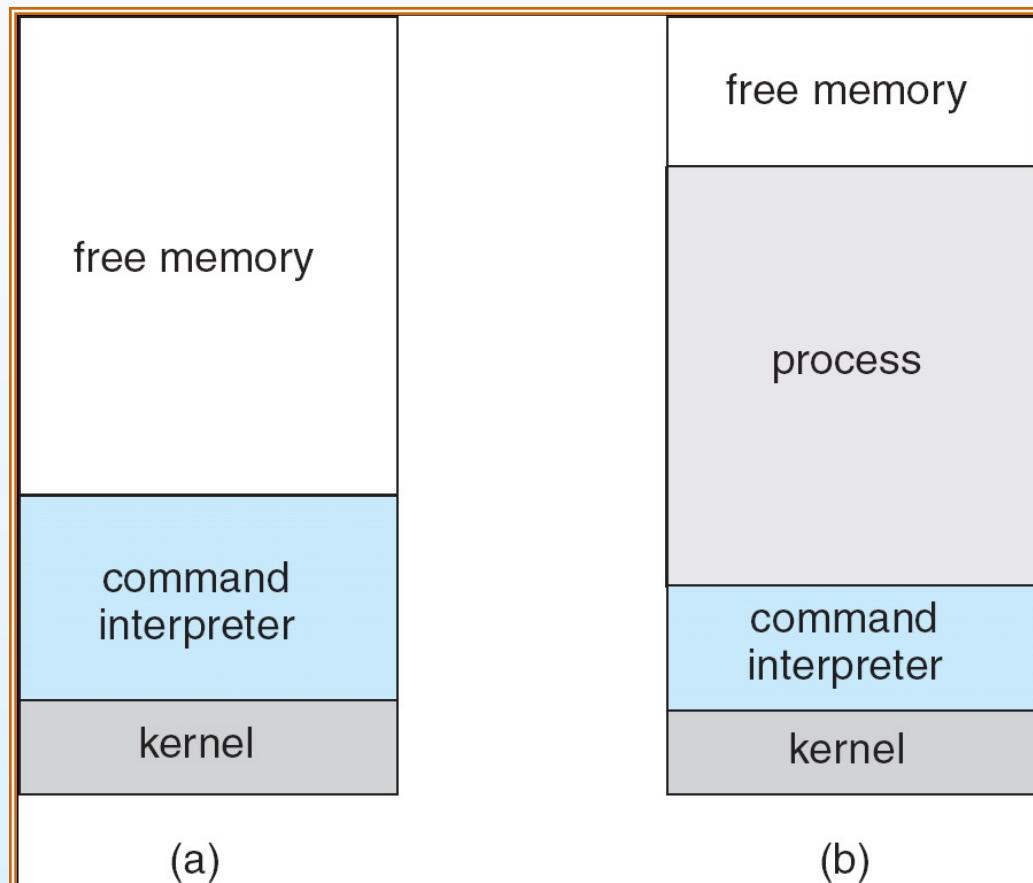
# Types of System Calls

- Process control
- File management
- Device management
- Information maintenance (e.g. time, date)
- Communications





# Example of process control: MS-DOS execution

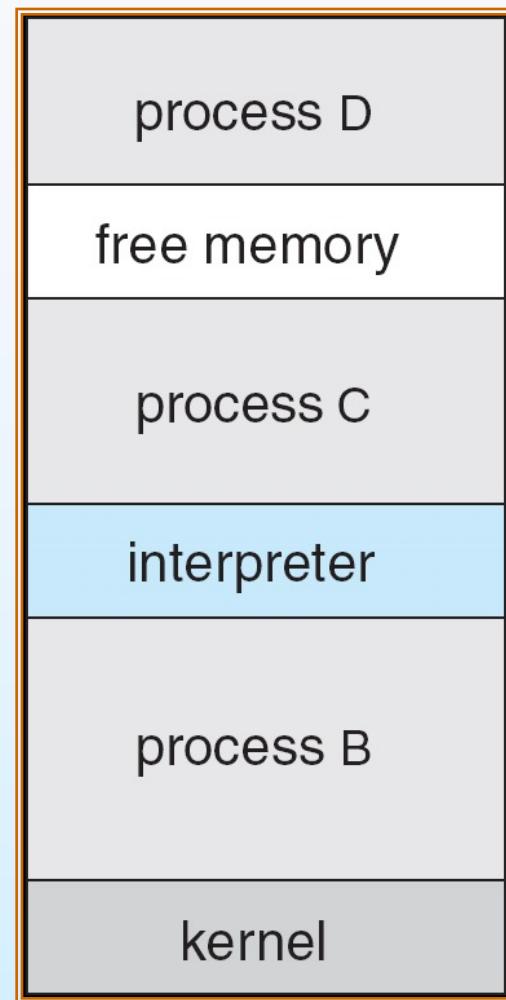


(a) At system startup (b) running a  
program, overwriting the OS itself





# FreeBSD Running Multiple Programs





# Chapter 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- System Calls
- Types of System Calls
- **System Programs**
- Operating System Design and Implementation
- Operating System Structure
- Virtual Machines
- Operating System Generation
- System Boot

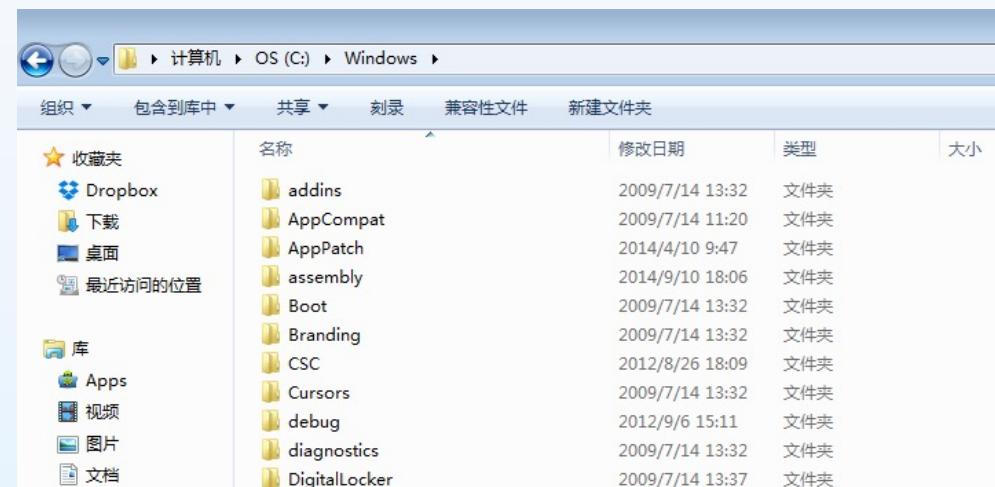




# System Programs

- System programs provide a convenient environment for program development and execution. They can be divided into:

- File manipulation
- Status information
- File modification
- Communications
- Application programs
- Programming language support
- Program loading and execution



名称	修改日期	类型
addons	2009/7/14 13:32	文件夹
AppCompat	2009/7/14 11:20	文件夹
AppPatch	2014/4/10 9:47	文件夹
assembly	2014/9/10 18:06	文件夹
Boot	2009/7/14 13:32	文件夹
Branding	2009/7/14 13:32	文件夹
CSC	2012/8/26 18:09	文件夹
Cursors	2009/7/14 13:32	文件夹
debug	2012/9/6 15:11	文件夹
diagnostics	2009/7/14 13:32	文件夹
DigitalLocker	2009/7/14 13:37	文件夹

- Most users' view of the operating system is defined by system programs, not the actual system calls





# Solaris 10 dtrace Following System Call

```
# ./all.d `pgrep xclock` XEventsQueued
dtrace: script './all.d' matched 52377 probes
CPU FUNCTION
  0 -> XEventsQueued                      U
  0  -> _XEventsQueued                     U
  0  -> _X11TransBytesReadable              U
  0  <- _X11TransBytesReadable              U
  0  -> _X11TransSocketBytesReadable       U
  0  <- _X11TransSocketBytesreadable        U
  0  -> ioctl                            U
  0  -> ioctl                            K
  0  -> getf                             K
  0    -> set_active_fd                  K
  0    <- set_active_fd                  K
  0  <- getf                            K
  0  -> get_udatamodel                 K
  0  <- get_udatamodel                 K
...
  0  -> releaseef                      K
  0    -> clear_active_fd               K
  0    <- clear_active_fd               K
  0    -> cv_broadcast                  K
  0    <- cv_broadcast                  K
  0    <- releaseef                    K
  0    <- ioctl                        K
  0  <- ioctl                        U
  0  <- _XEventsQueued                U
  0 <- XEventsQueued                  U
```





# System Programs

- Provide a **convenient environment** for program development and execution
  - Some of them are simply user interfaces to system calls; others are considerably more complex
- **File management** - Create, delete, copy, rename, print, dump, list, and generally manipulate files and directories
- **Status information**
  - Some ask the system for info - date, time, amount of available memory, disk space, number of users
  - Others provide detailed performance, logging, and debugging information
  - Typically, these programs format and print the output to the terminal or other output devices
  - Some systems implement a registry - used to store and retrieve configuration information





# System Programs (cont'd)

- **File modification**
  - Text editors to create and modify files
  - Special commands to search contents of files or perform transformations of the text
- **Programming-language support** - Compilers, assemblers, debuggers and interpreters sometimes provided
- **Program loading and execution**- Absolute loaders, relocatable loaders, linkage editors, and overlay-loaders, debugging systems for higher-level and machine language
- **Communications** - Provide the mechanism for creating virtual connections among processes, users, and computer systems
  - Allow users to send messages to one another's screens, browse web pages, send electronic-mail messages, log in remotely, transfer files from one machine to another





# Chapter 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- System Calls
- Types of System Calls
- System Programs
- Operating System Design and Implementation
- Operating System Structure
- Virtual Machines
- Operating System Generation
- System Boot





# Operating System Design and Implementation

- Design and Implementation of OS **not “solvable”**, but some approaches have proven successful
- Internal structure of different Operating Systems can vary widely
- Start by defining **goals** and **specifications**
- Affected by choice of **hardware**, **type of system**
- *User goals and System goals*
  - User goals – operating system should be convenient to use, easy to learn, reliable, safe, and fast
  - System goals – operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient
- Successful design: Windows 95, Windows XP, OSX
- Failed design: Windows vista, Windows 8, Android 1.x

Creative process





# Operating System Design and Implementation (Cont.)

- Important principle to separate (2.6.2)

**Policy:** What will be done?

**Mechanism:** How to do it?

Example: timer for CPU protection is a mechanism

- Mechanisms determine how to do something, policies decide what will be done
  - The separation of policy from mechanism is a very important principle, it allows maximum flexibility if policy decisions are to be changed later





## Example: Separation of Mechanism and Policy

- An example of mechanism/policy separation is the use of **card-keys** to gain access to locked doors.
- ***The mechanisms*** (magnetic card readers, remote controlled locks, connections to a security server)
- ***The entrance policy*** (which people should be allowed to enter which doors, at which times).
- These decisions are made by a centralized security server, which (in turn) probably makes its decisions by consulting a database of room access rules.





# Additional readings: The early principles of X

A very good illustration of the separation: X-Window System

- **Do not add new functionality unless an implementor cannot complete a real application without it.**
- **It is as important to decide what a system is not as to decide what it is. Do not serve all the world's needs; rather, make the system extensible so that additional needs can be met in an upwardly compatible fashion. (windows 8)**
- **The only thing worse than generalizing from one example is generalizing from no examples at all.**
- **If a problem is not completely understood, it is probably best to provide no solution at all.**
- **If you can get 90 percent of the desired effect for 10 percent of the work, use the simpler solution.**
- **Isolate complexity as much as possible.**
- **Provide mechanism rather than policy. In particular, place user interface policy in the clients' hands.**





# Chapter 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- System Calls
- Types of System Calls
- System Programs
- Operating System Design and Implementation
- **Operating System Structure**
- Virtual Machines
- Operating System Generation
- System Boot





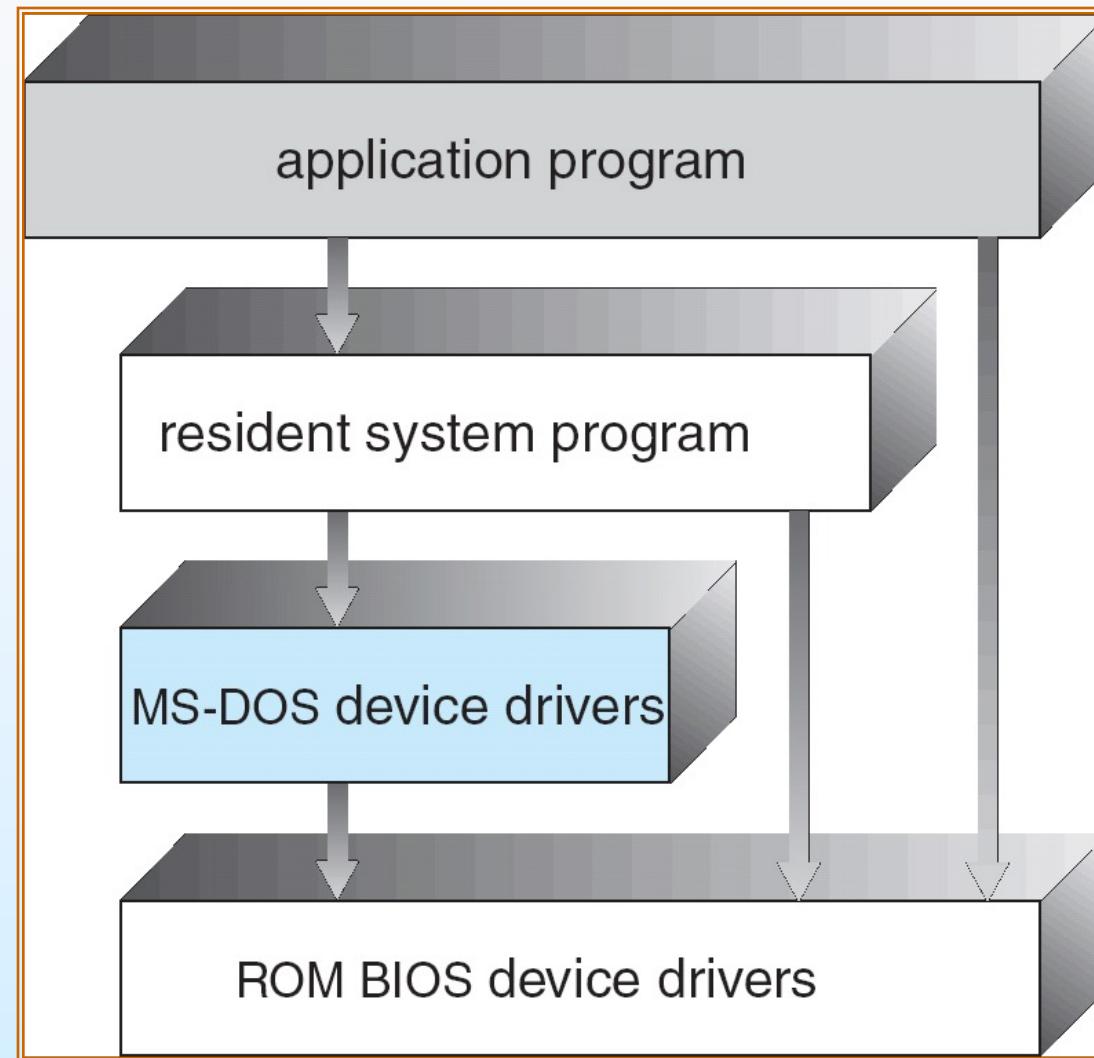
# Simple Structure

- MS-DOS – written to provide the most functionality in the least space
  - Not divided into modules
  - Although MS-DOS has some structure, its interfaces and levels of functionality are not well separated





# MS-DOS Layer Structure





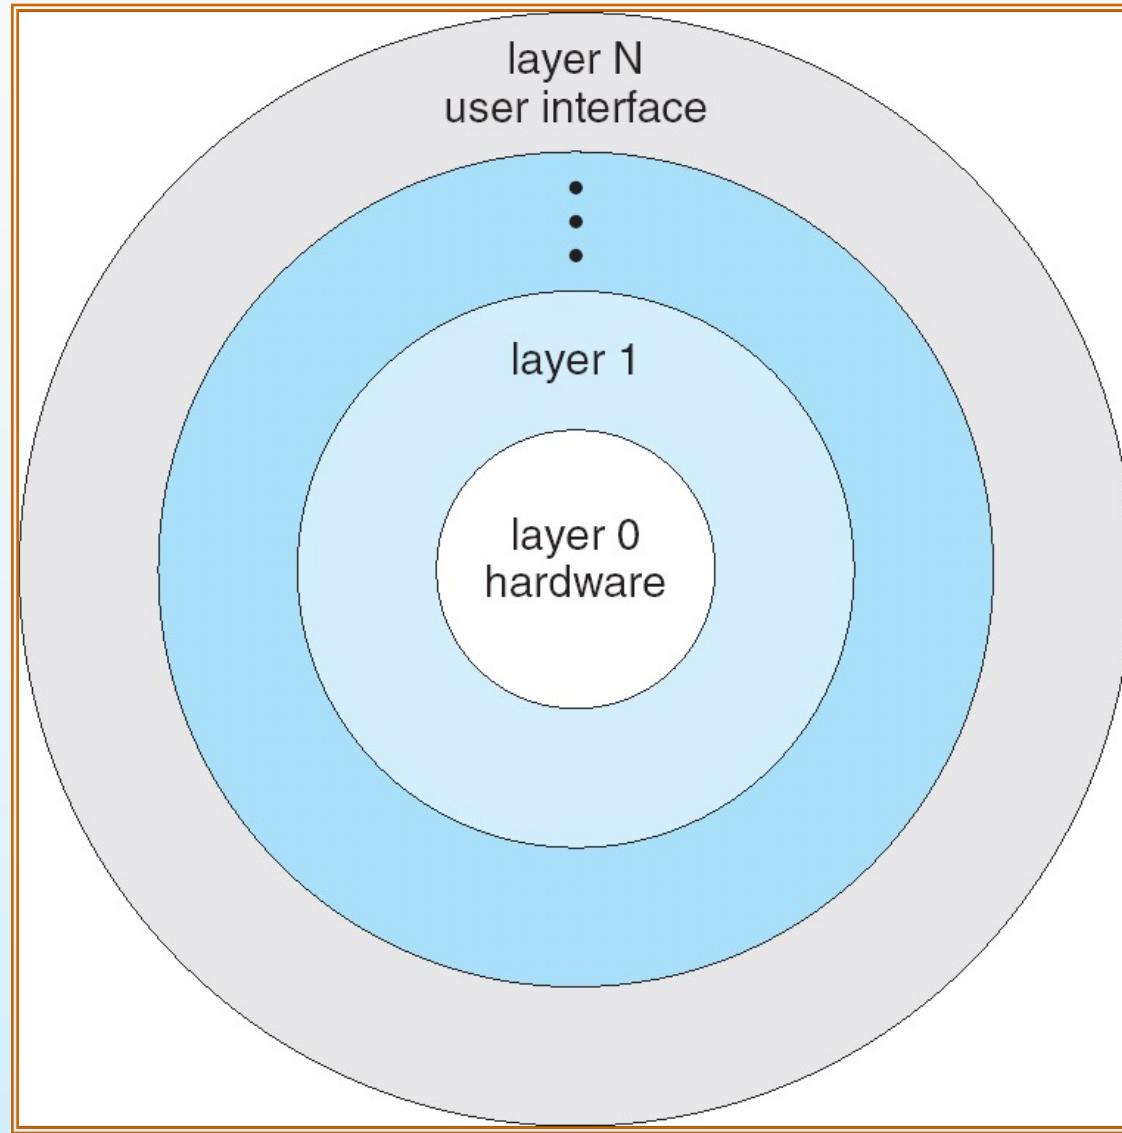
# Layered Approach

- The operating system is divided into a number of layers (levels), each built on top of lower layers. The bottom layer (layer 0), is the hardware; the highest (layer N) is the user interface.
- With modularity, layers are selected such that each uses functions (operations) and services of only lower-level layers





# Layered Operating System





# UNIX

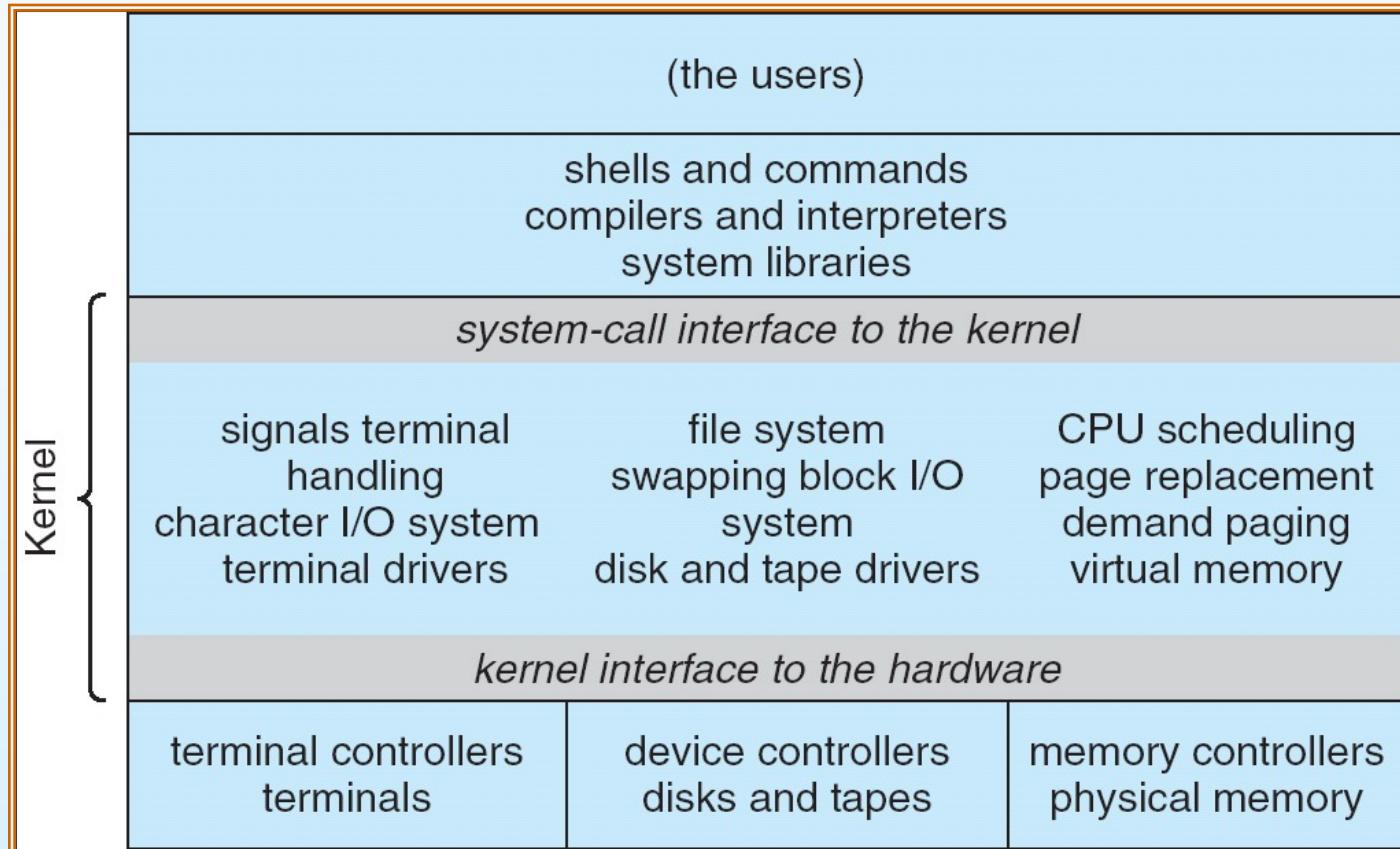
- UNIX – limited by hardware functionality, the original UNIX operating system had limited structuring. The UNIX OS consists of two separable parts
  - System programs
  - The kernel
    - ▶ Consists of everything below the system-call interface and above the physical hardware
    - ▶ Provides the file system, CPU scheduling, memory management, and other operating-system functions; a large number of functions for one level

**Monolithic structure!**





# UNIX System Structure





# Microkernel System Structure

- Moves as much from the kernel into “*user*” space
- Communication takes place between **user modules** using **message passing**
- Benefits:
  - Easier to **extend** a microkernel
  - Easier to **port** the operating system to new architectures
  - More **reliable** (less code is running in kernel mode)
  - More **secure**
- Detriments:
  - **Performance overhead** of user space to kernel space communication





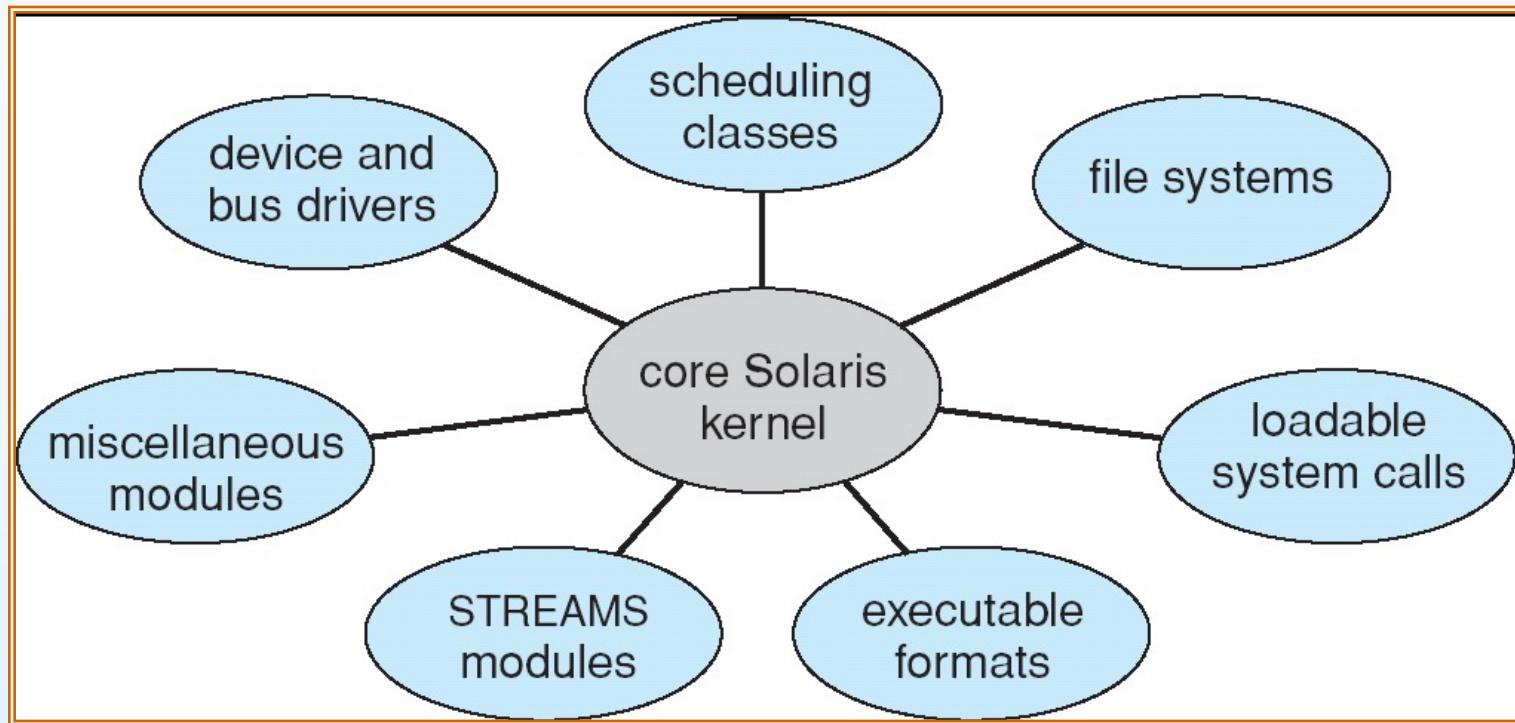
# Modules

- Most modern operating systems implement **kernel modules**
  - Uses object-oriented approach
  - Each core component is separate
  - Each talks to the others over known interfaces
    - ▶ Modules call each other instead of message passing
  - Each is **loadable** as needed within the kernel
- Overall, similar to layers but more flexible





# Solaris Modular Approach

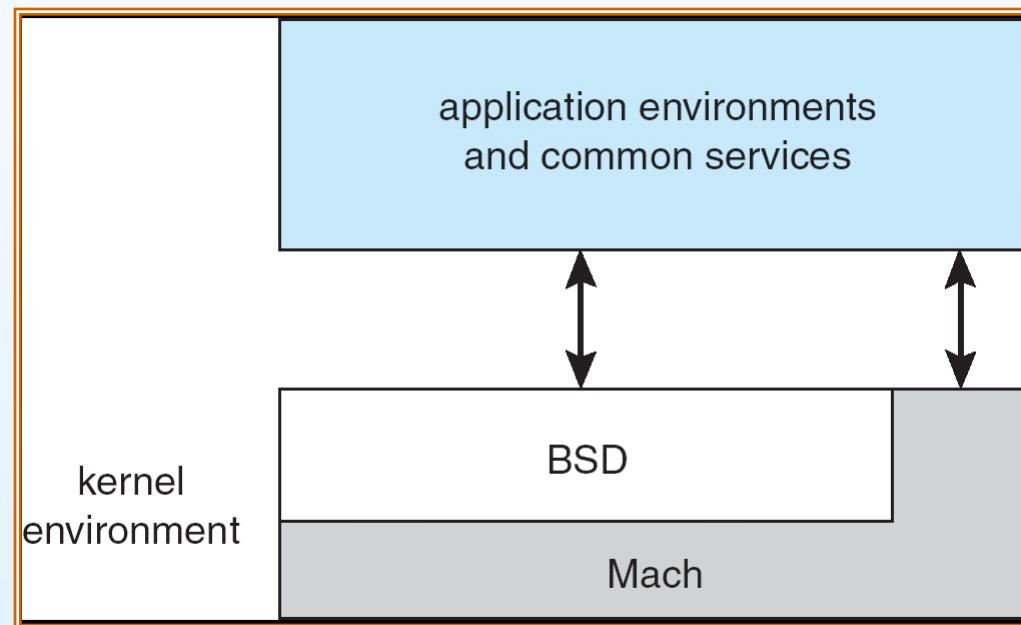


**Modular kernel**





# Mac OS X Structure



A hybrid structure!





# Chapter 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- System Calls
- Types of System Calls
- System Programs
- Operating System Design and Implementation
- Operating System Structure
- **Virtual Machines**
- Operating System Generation
- System Boot





# Virtual Machines

- A *virtual machine* takes the layered approach to its logical conclusion. It treats hardware and the operating system kernel as though they were all **hardware**
- A virtual machine provides an interface *identical* to the underlying bare hardware
- The operating system creates the illusion of multiple processes, each executing on its own processor with its own (virtual) memory





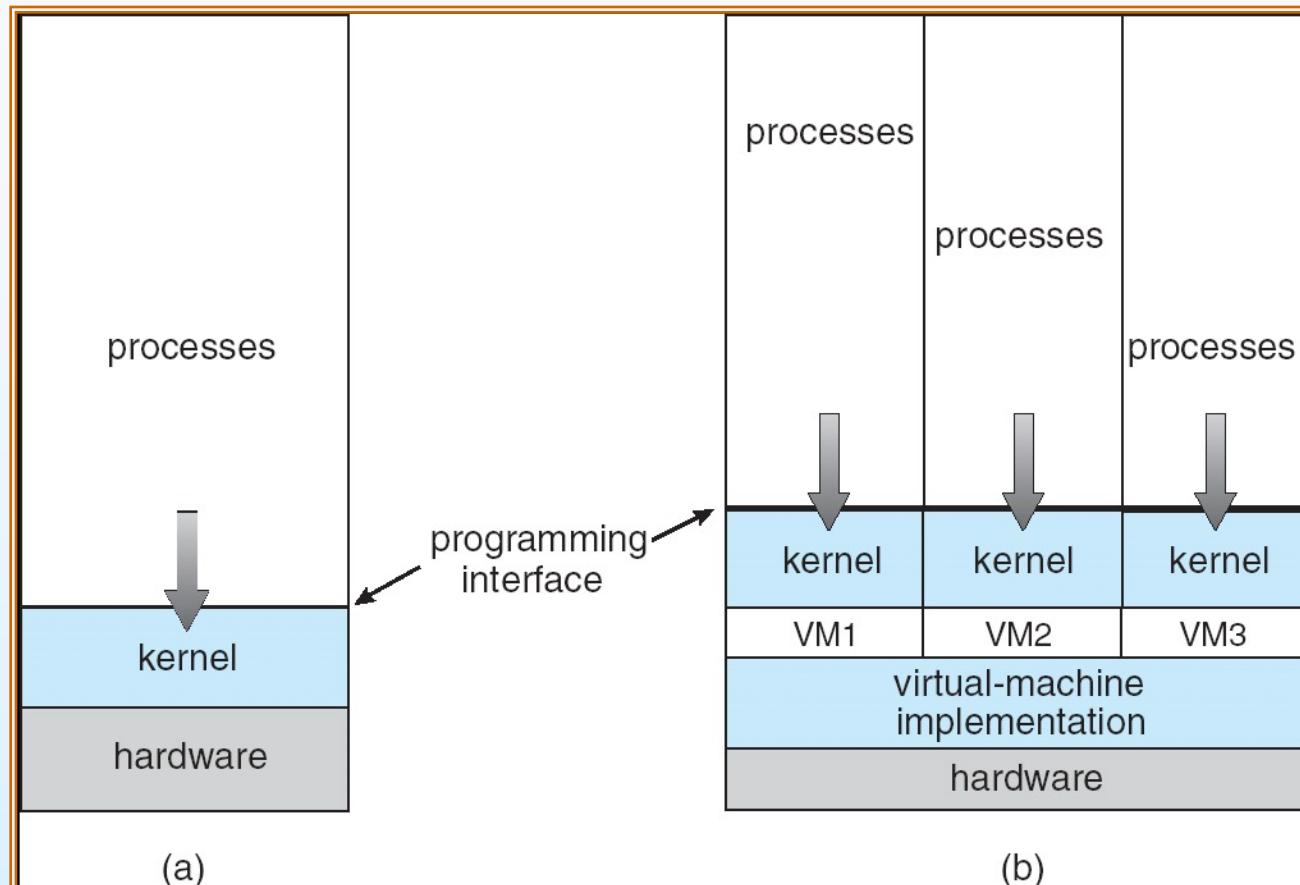
# Virtual Machines (Cont.)

- The resources of the physical computer are shared to create the virtual machines
  - CPU scheduling can create the appearance that users have their own processor
  - Spooling and a file system can provide virtual card readers and virtual line printers
  - A normal user time-sharing terminal serves as the virtual machine operator's console





# Virtual Machines (Cont.)



(a) Nonvirtual machine

(b) virtual machine





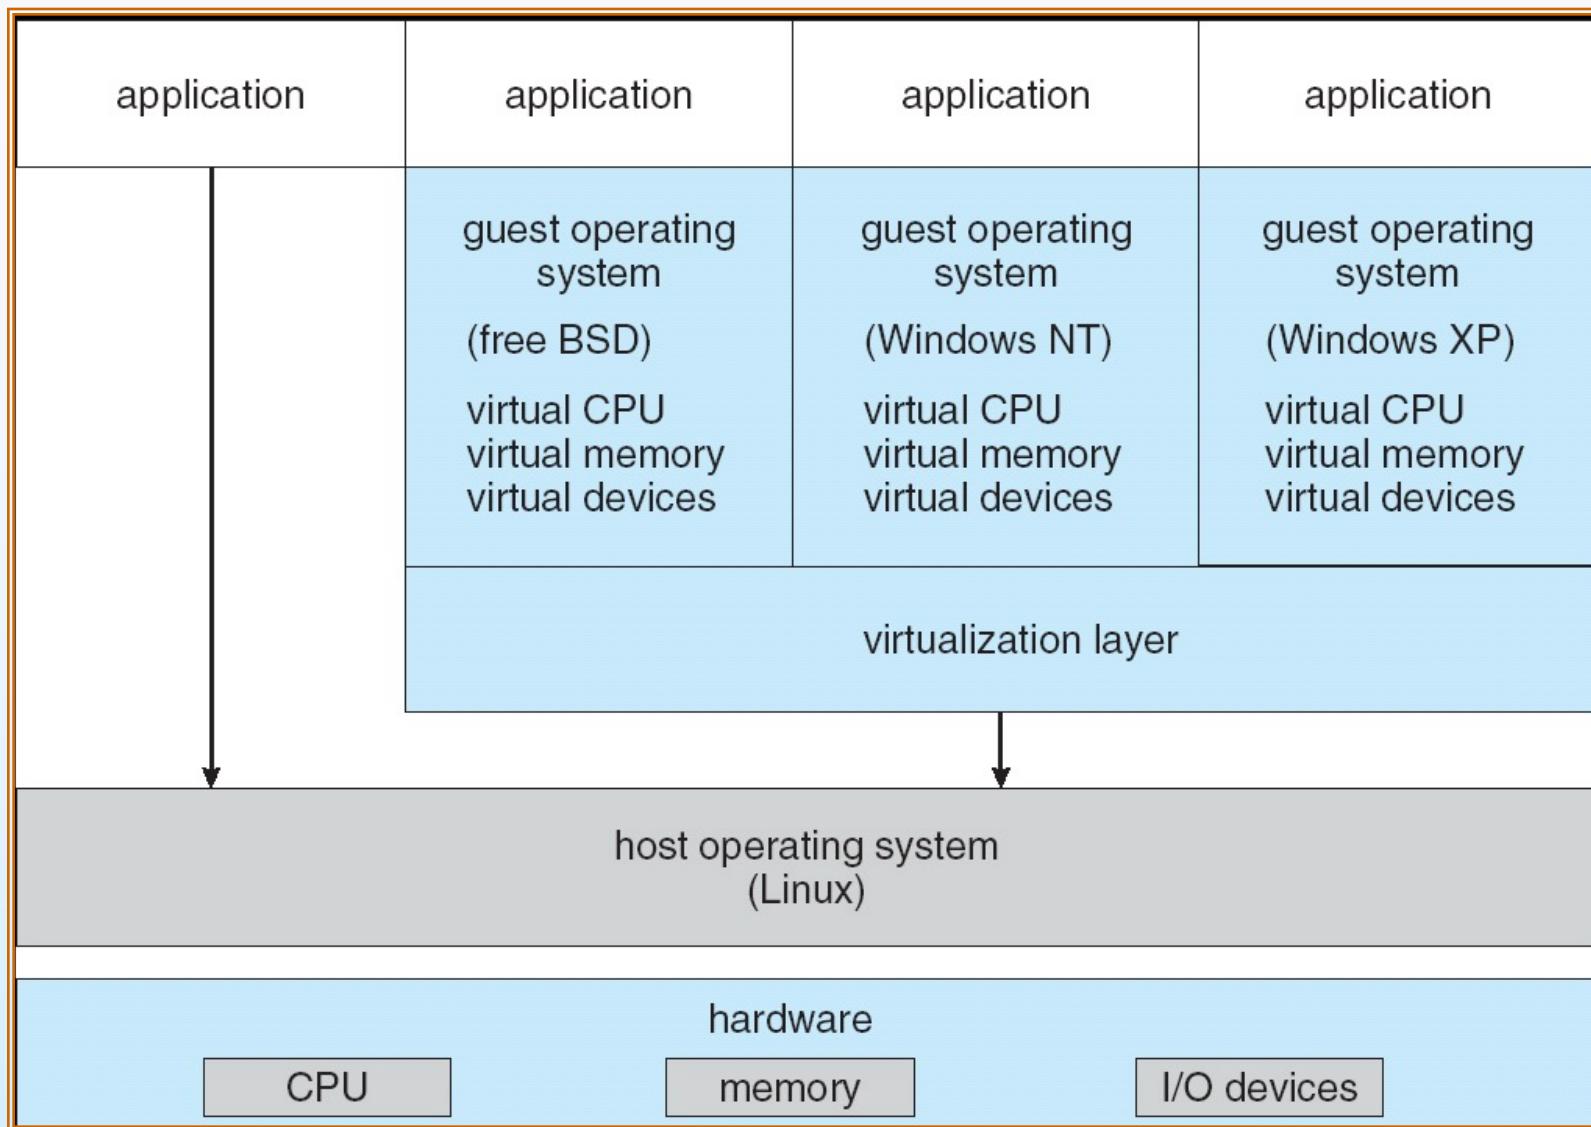
# Virtual Machines (Cont.)

- The virtual-machine concept provides complete protection of system resources since each virtual machine is isolated from all other virtual machines. This isolation, however, permits no direct sharing of resources.
- A virtual-machine system is a perfect vehicle for operating-systems research and development. **System development** is done on the virtual machine, instead of on a physical machine and so does not disrupt normal system operation.
- The virtual machine concept is difficult to implement due to the effort required to provide an ***exact*** duplicate to the underlying machine. (for example, virtual user mode and kernel mode)





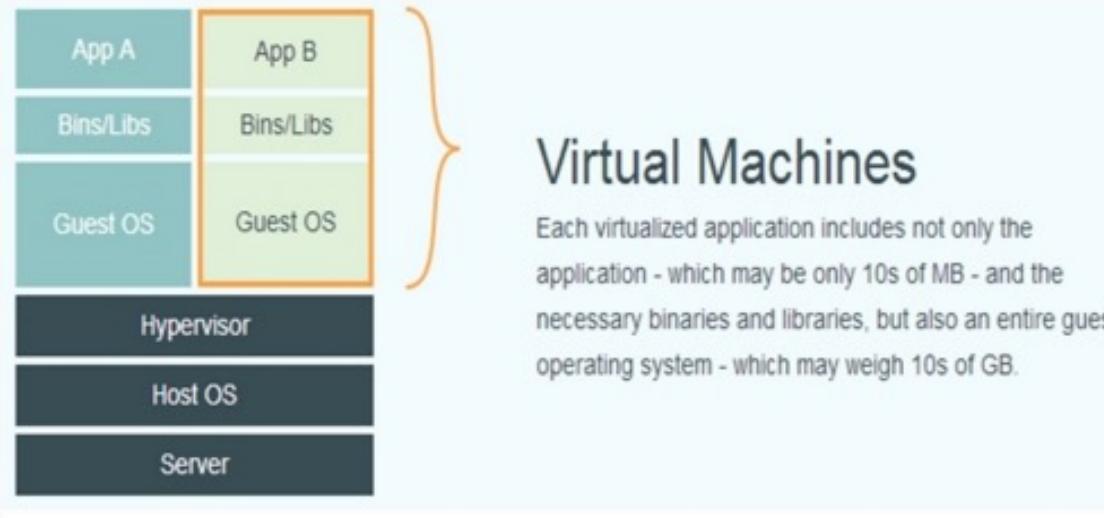
# VMware Architecture



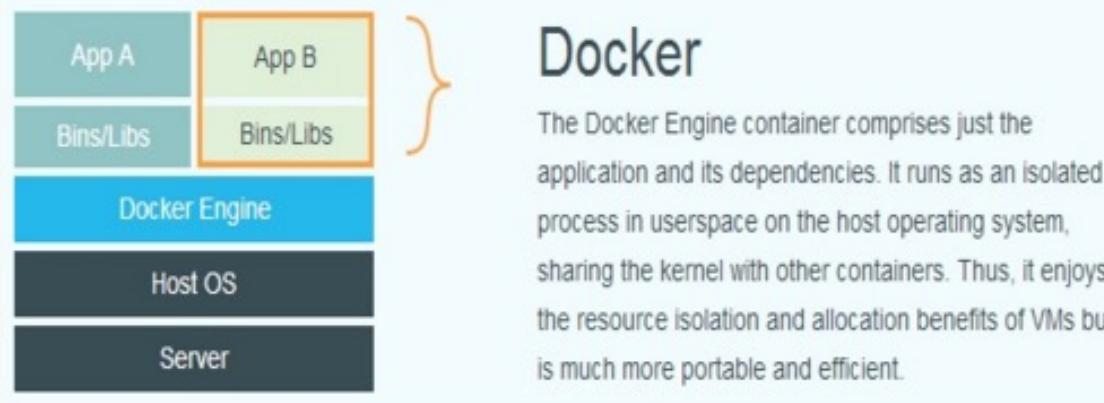


# VM vs Docker

下图为传统虚拟化方案：

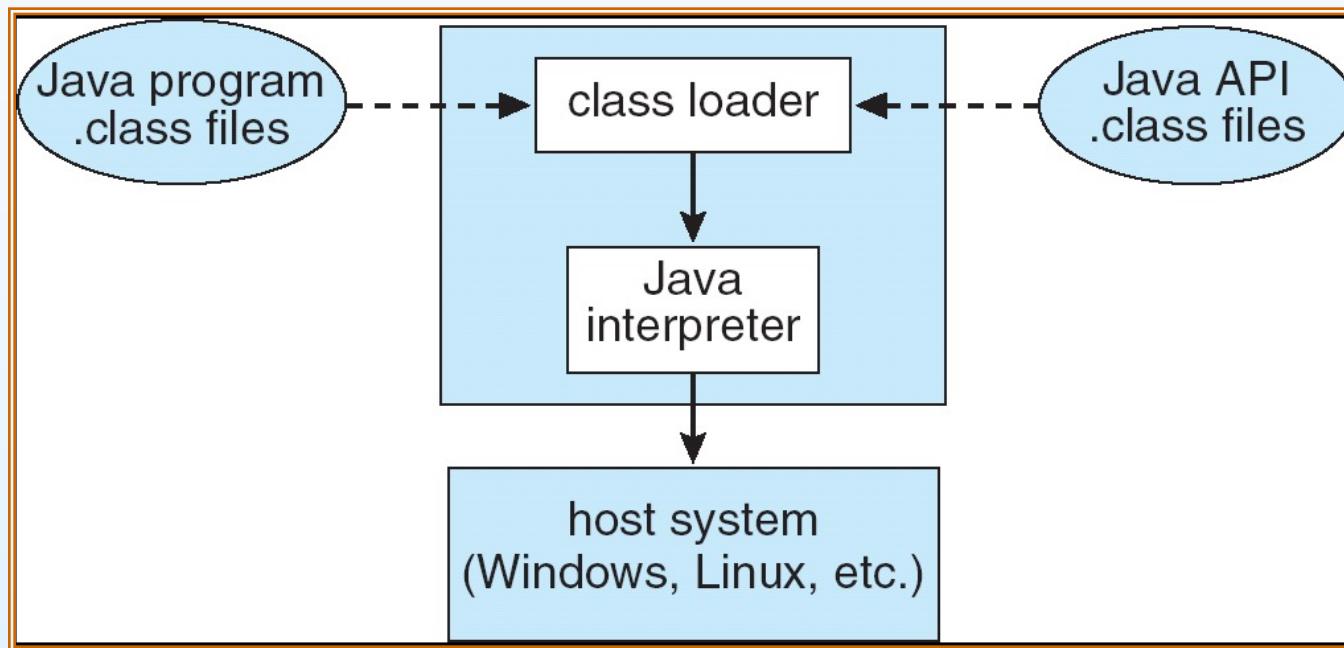


如下为Docker虚拟化方案：





# The Java Virtual Machine





# Chapter 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- System Calls
- Types of System Calls
- System Programs
- Operating System Design and Implementation
- Operating System Structure
- Virtual Machines
- **Operating System Generation**
- System Boot





# Operating System Generation

- Operating systems are designed to run on any of a class of machines; the system must be configured for each specific computer site
- SYSGEN program obtains information concerning the specific configuration of the hardware system
- *Booting* – starting a computer by loading the kernel
- *Bootstrap program* – code stored in ROM that is able to locate the kernel, load it into memory, and start its execution





# Chapter 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- System Calls
- Types of System Calls
- System Programs
- Operating System Design and Implementation
- Operating System Structure
- Virtual Machines
- Operating System Generation
- System Boot





# System Boot

- Operating system must be made available to hardware so hardware can start it
  - Small piece of code – **bootstrap program (a.k.a. bootstrap loader)**, locates the kernel, loads it into memory, and starts it
  - Sometimes two-step process where **boot block** at fixed location loads bootstrap loader
  - When power initialized on system, execution starts at a fixed memory location
    - ▶ Firmware used to hold initial boot code





The picture can't be displayed.

## End of Chapter 2

